

```
from google.colab import files
uploaded = files.upload()
```

```
<IPython.core.display.HTML object>
```

```
Saving Stores.csv to Stores (1).csv
```

```
import pandas as pd
```

```
df = pd.read_csv("Stores.csv")
```

```
print(df.head())
```

```
print(df.shape)
```

```
print(df.columns)
```

	Store ID	Store_Area	Items_Available	Daily_Customer_Count	Store_Sales
0	1	1659	1961	530	66490
1	2	1461	1752	210	39820
2	3	1340	1609	720	54010
3	4	1451	1748	620	53730
4	5	1770	2111	450	46620

(896, 5)

```
Index(['Store ID ', 'Store_Area', 'Items_Available',  
      'Daily_Customer_Count',  
      'Store_Sales'],  
      dtype='object')
```

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 896 entries, 0 to 895
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	Store ID	896 non-null	int64
1	Store_Area	896 non-null	int64
2	Items_Available	896 non-null	int64
3	Daily_Customer_Count	896 non-null	int64
4	Store_Sales	896 non-null	int64

```
dtypes: int64(5)
```

```
memory usage: 35.1 KB
```

```
None
```

```
print(df.describe(include="all").T)
```

	count	mean	std	min
25% \				
Store_ID	896.0	448.500000	258.797218	1.0
224.75				
Store_Area	896.0	1485.409598	250.237011	775.0
1316.75				
Items_Available	896.0	1782.035714	299.872053	932.0
1575.50				
Daily_Customer_Count	896.0	786.350446	265.389281	10.0
600.00				
Store_Sales	896.0	59351.305804	17190.741895	14920.0
46530.00				

	50%	75%	max
Store_ID	448.5	672.25	896.0
Store_Area	1477.0	1653.50	2229.0
Items_Available	1773.5	1982.75	2667.0
Daily_Customer_Count	780.0	970.00	1560.0
Store_Sales	58605.0	71872.50	116320.0

```
print(df.isnull().sum())
```

```
Store_ID      0
Store_Area    0
Items_Available  0
Daily_Customer_Count  0
Store_Sales    0
dtype: int64
```

```
for col in df.columns:
    print(f"{col} -> {df[col].nunique()} unique values")
```

```
Store_ID -> 896 unique values
Store_Area -> 583 unique values
Items_Available -> 616 unique values
Daily_Customer_Count -> 130 unique values
Store_Sales -> 816 unique values
```

```
print(df.skew())
```

```
Store_ID      0.000000
Store_Area    0.030367
Items_Available  0.034439
Daily_Customer_Count  0.074633
Store_Sales    0.148794
dtype: float64
```

```
for col in df.select_dtypes(include=['int64','float64']).columns:
    Q1 = df[col].quantile(0.25)
```


```
Q3 = df[col].quantile(0.75)
IQR = Q3 - Q1
lower = Q1 - 1.5*IQR
upper = Q3 + 1.5*IQR
print(f"{col}: {((df[col] < lower) | (df[col] > upper)).sum()}
outliers")
```

```
Store_ID : 0 outliers
Store_Area: 5 outliers
Items_Available: 5 outliers
Daily_Customer_Count: 3 outliers
Store_Sales: 1 outliers
```

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns


df = pd.read_csv("Stores.csv")

sns.heatmap(df.isnull(), cbar=False, cmap="viridis")
plt.title("Missing Values Heatmap")
plt.show()
```





```
for col in df.columns[1:]:  
    sns.histplot(df[col], kde=True, bins=30)  
    plt.title(f"Histogram of {col}")  
    plt.show()
```


Histogram of Store_Area



Histogram of Items_Available









```
for col in df.columns[1:]:  
    sns.boxplot(x=df[col])  
    plt.title(f"Boxplot of {col}")  
    plt.show()
```



Boxplot of Store_Area

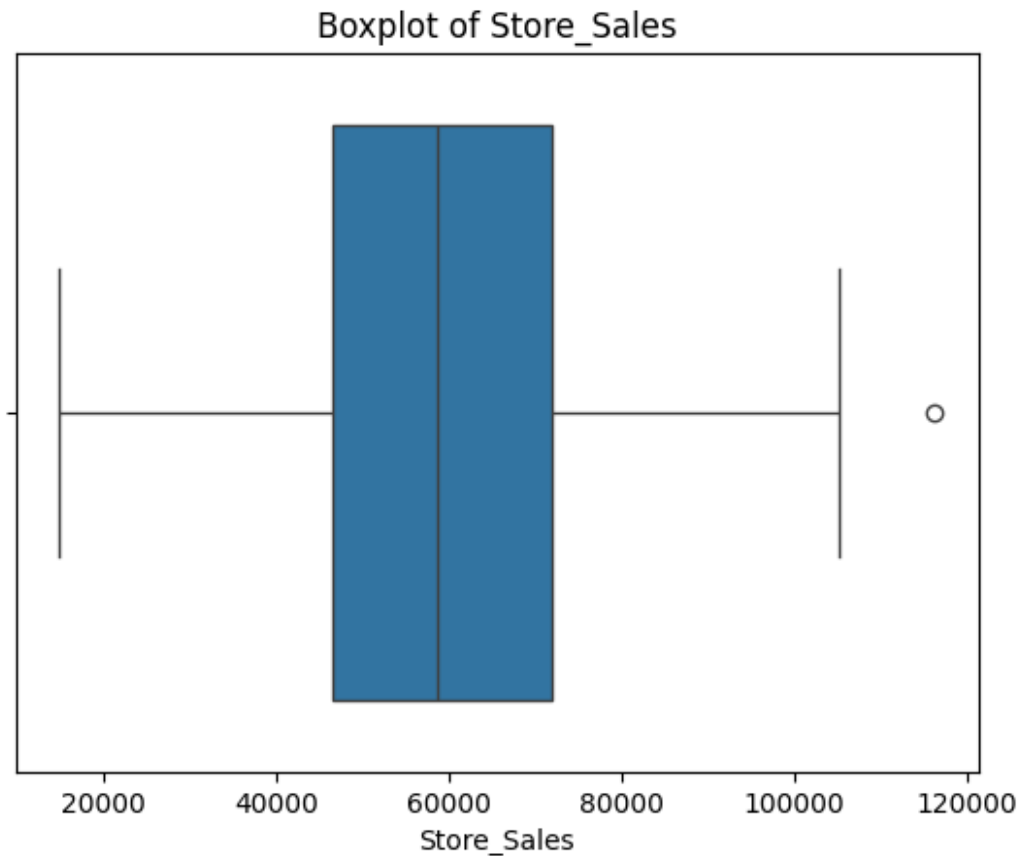


Boxplot of Items_Available



Boxplot of Daily_Customer_Count



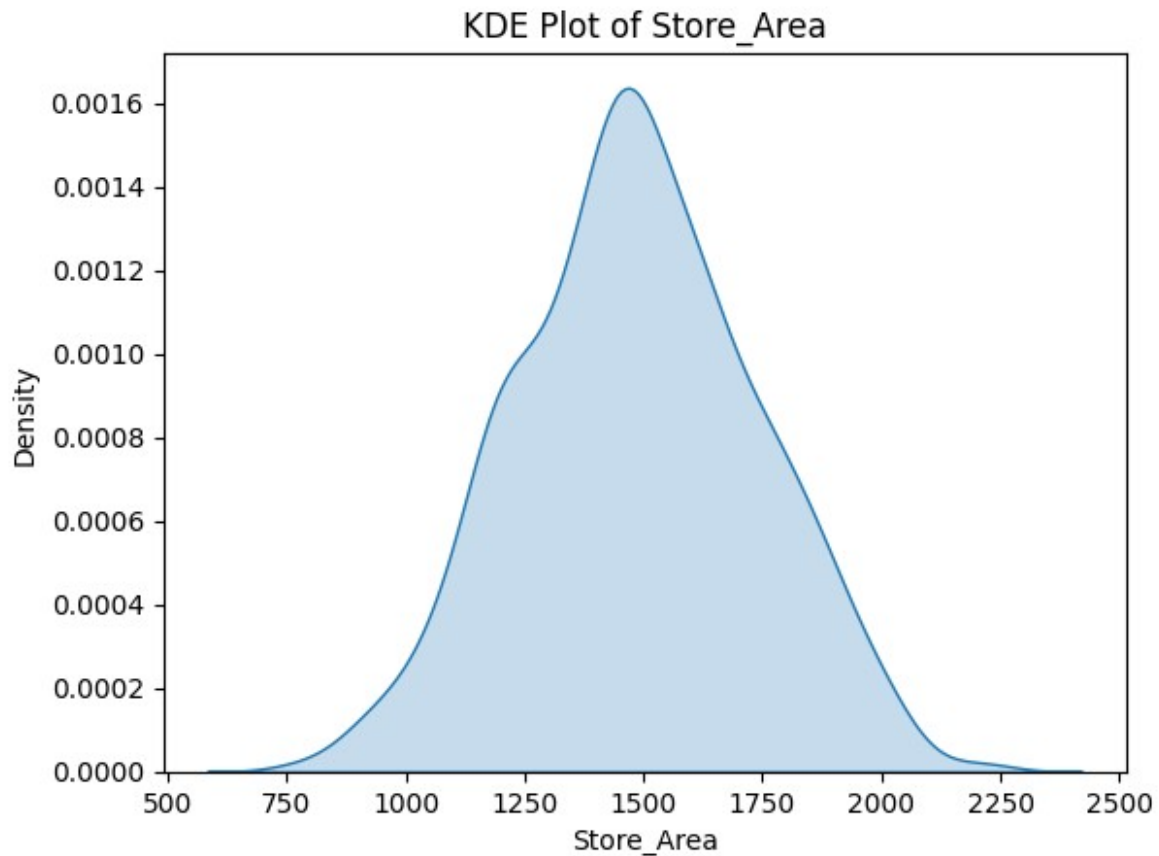


```
for col in df.columns[1:]:  
    sns.kdeplot(df[col], shade=True)  
    plt.title(f"KDE Plot of {col}")  
    plt.show()
```

/tmp/ipython-input-1999266753.py:2: FutureWarning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

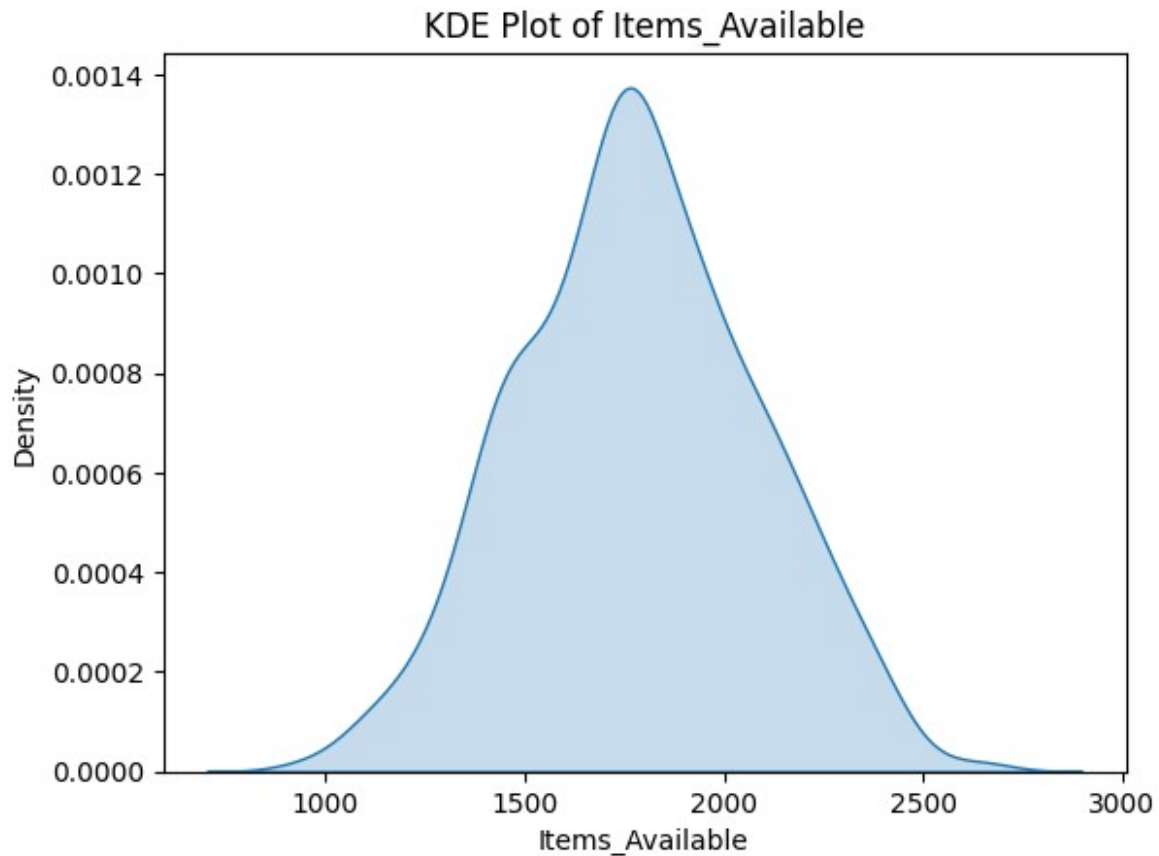
```
sns.kdeplot(df[col], shade=True)
```



```
/tmp/ipython-input-1999266753.py:2: FutureWarning:
```

```
`shade` is now deprecated in favor of `fill`; setting `fill=True`.  
This will become an error in seaborn v0.14.0; please update your code.
```


```
sns.kdeplot(df[col], shade=True)
```



```
/tmp/ipython-input-1999266753.py:2: FutureWarning:
```

```
`shade` is now deprecated in favor of `fill`; setting `fill=True`.  
This will become an error in seaborn v0.14.0; please update your code.
```


```
sns.kdeplot(df[col], shade=True)
```



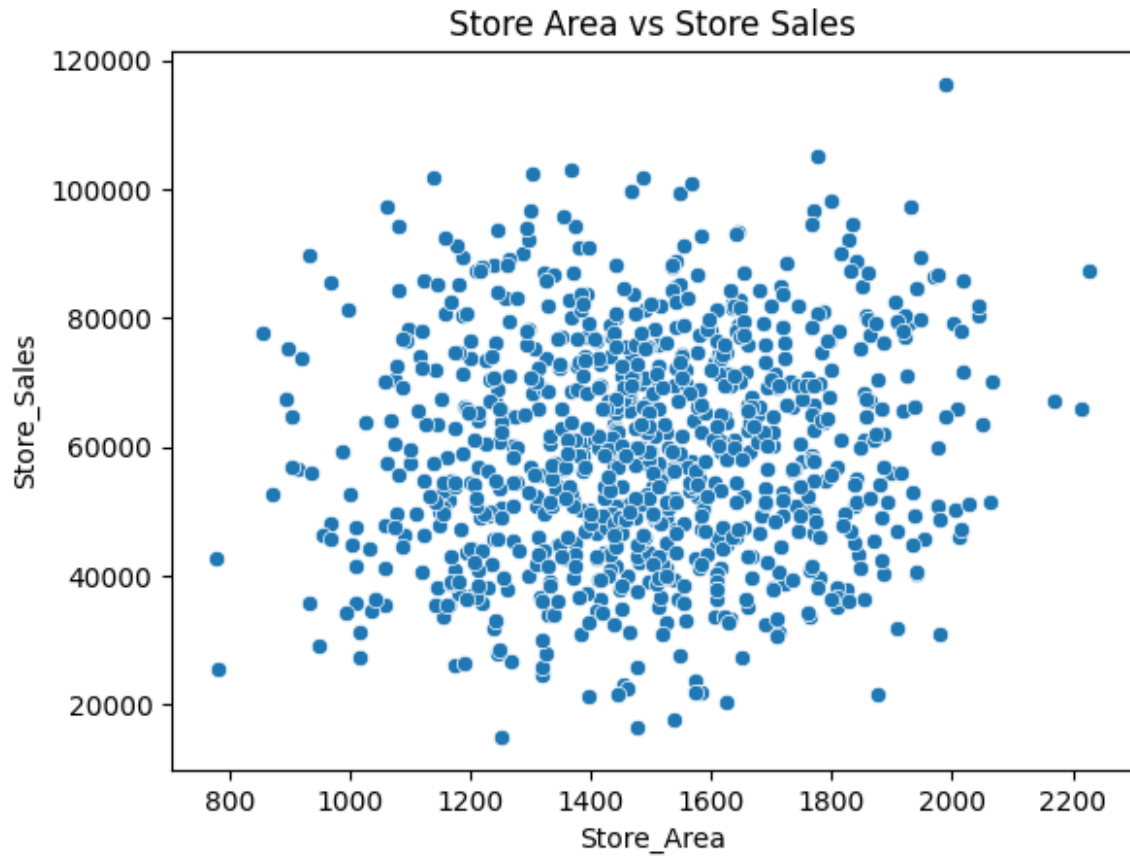
```
/tmp/ipython-input-1999266753.py:2: FutureWarning:
```

```
`shade` is now deprecated in favor of `fill`; setting `fill=True`.  
This will become an error in seaborn v0.14.0; please update your code.
```

```
sns.kdeplot(df[col], shade=True)
```



```
sns.scatterplot(x="Store_Area", y="Store_Sales", data=df)
plt.title("Store Area vs Store Sales")
plt.show()
```

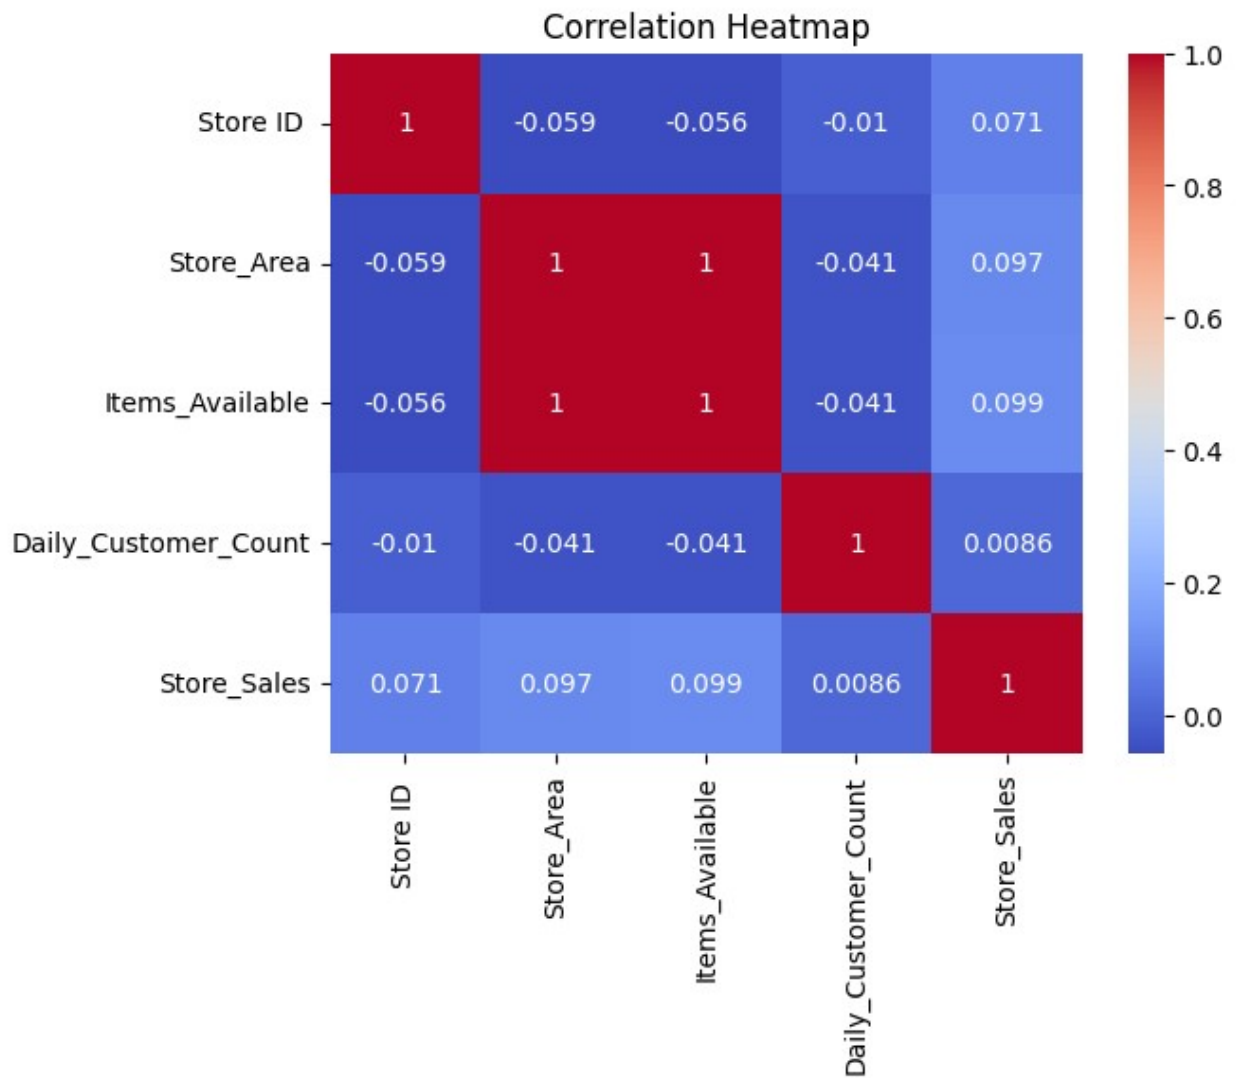
```
sns.scatterplot(x="Items_Available", y="Store_Sales", data=df)
plt.title("Items Available vs Store Sales")
plt.show()
```



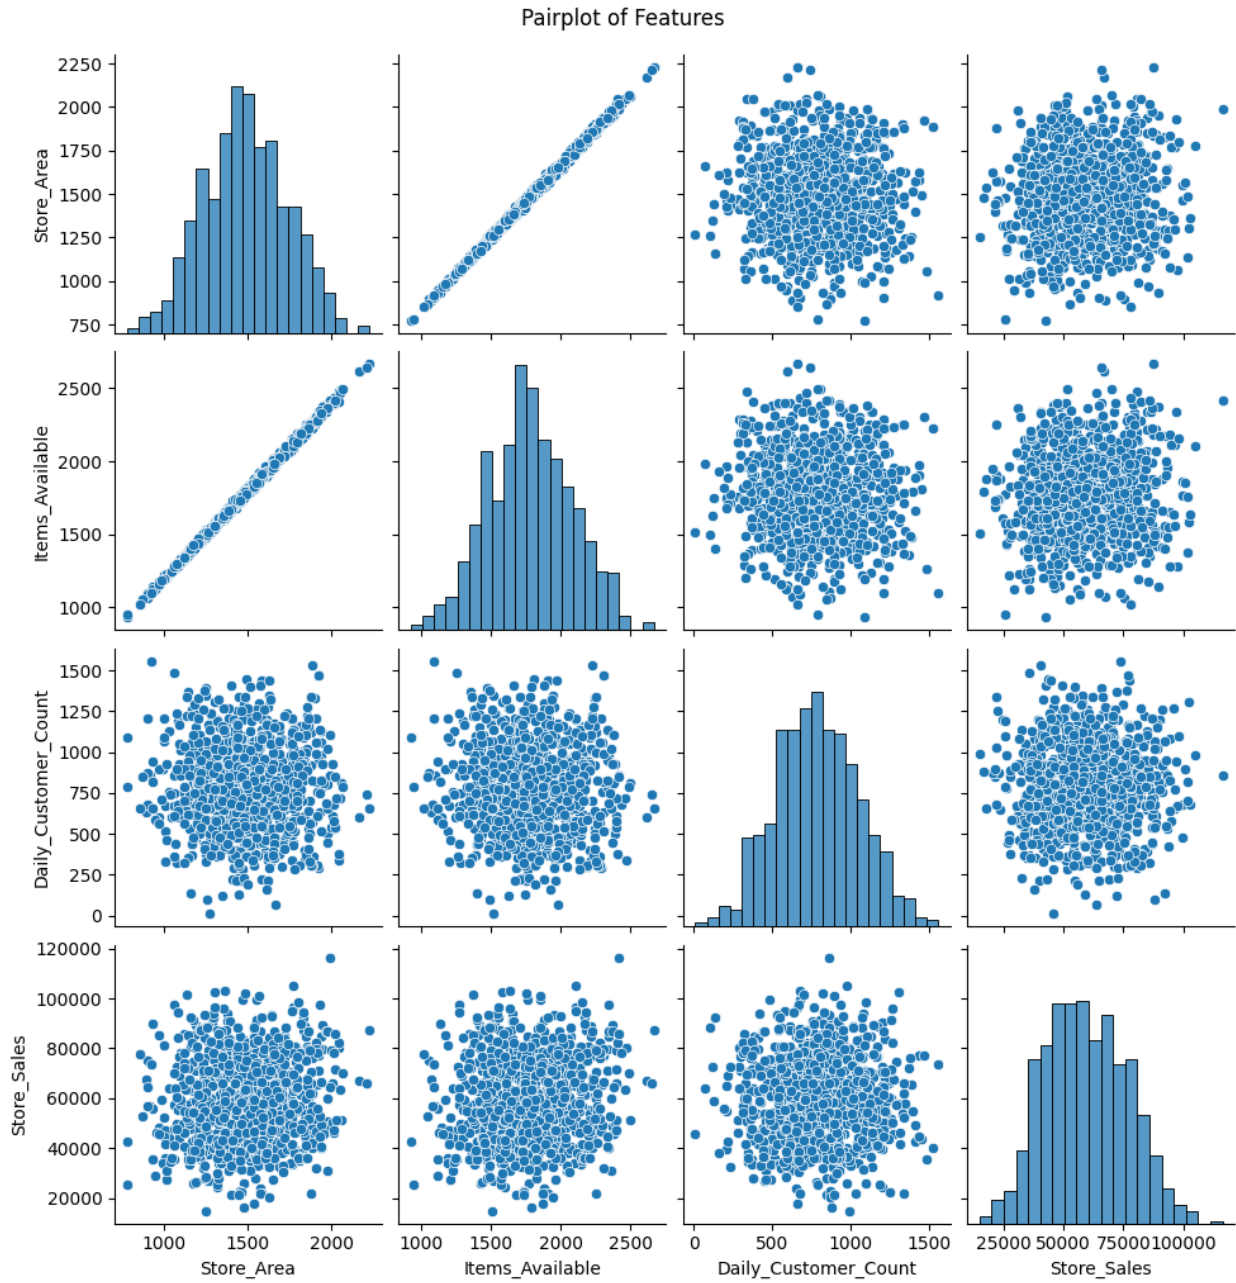
```
sns.scatterplot(x="Daily_Customer_Count", y="Store_Sales", data=df)
plt.title("Customers vs Store Sales")
plt.show()
```



```
sns.heatmap(df.corr(), annot=True, cmap="coolwarm")  
plt.title("Correlation Heatmap")  
plt.show()
```




```
sns.pairplot(df.drop("Store ID ", axis=1))  
plt.suptitle("Pairplot of Features", y=1.02)  
plt.show()
```




```
for col in ["Store_Area", "Items_Available", "Daily_Customer_Count",
"Store_Sales"]:
    sns.violinplot(x=df[col])
    plt.title(f"Violin Plot of {col}")
    plt.show()
```


Violin Plot of Store_Area

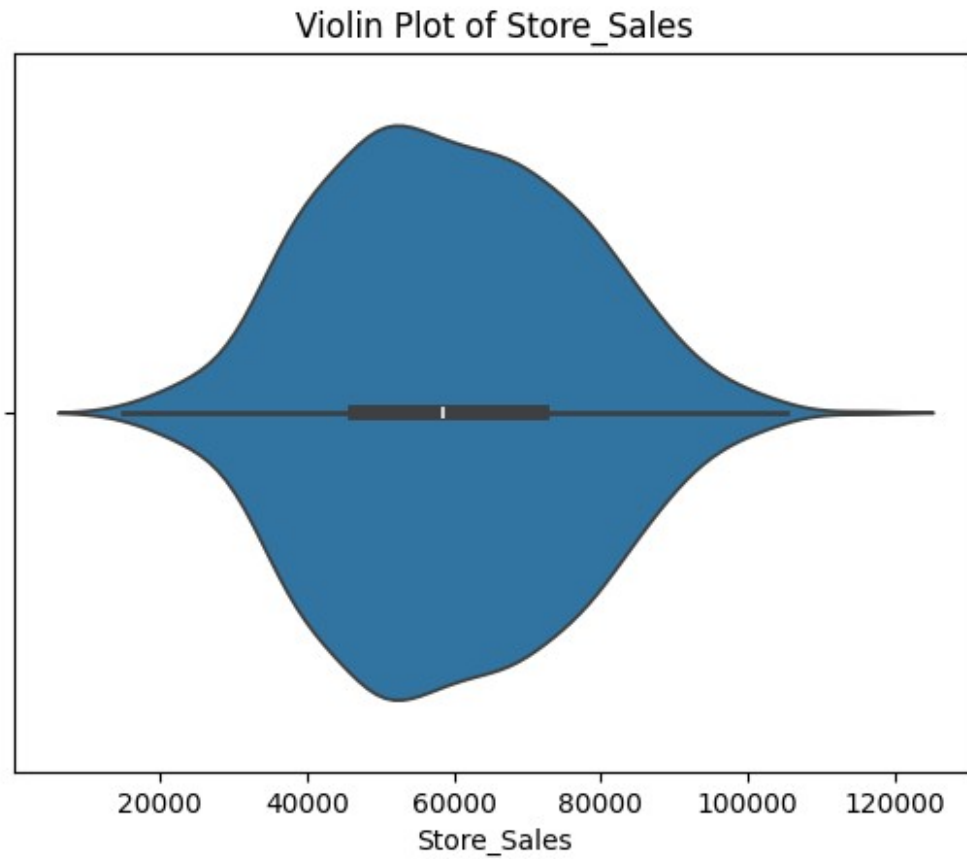


Violin Plot of Items_Available



Violin Plot of Daily_Customer_Count






```
sns.regplot(x="Store_Area", y="Store_Sales", data=df)
plt.title("Regression Plot: Store Area vs Store Sales")
plt.show()
```



```
sns.jointplot(x="Items_Available", y="Store_Sales", data=df,  
kind="hex")  
plt.show()
```



```
sns.swarmplot(x=df["Store_Area"], y=df["Store_Sales"])
plt.title("Swarmplot: Store Area vs Store Sales")
plt.show()
```

```
/usr/local/lib/python3.12/dist-packages/seaborn/categorical.py:3399:
UserWarning: 50.0% of the points cannot be placed; you may want to
decrease the size of the markers or use stripplot.
```


```
warnings.warn(msg, UserWarning)
```

```
/usr/local/lib/python3.12/dist-packages/seaborn/categorical.py:3399:
UserWarning: 33.3% of the points cannot be placed; you may want to
decrease the size of the markers or use stripplot.
```

```
warnings.warn(msg, UserWarning)
```

```
/usr/local/lib/python3.12/dist-packages/seaborn/categorical.py:3399:
UserWarning: 20.0% of the points cannot be placed; you may want to
decrease the size of the markers or use stripplot.
```


```
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.12/dist-packages/seaborn/categorical.py:3399:
UserWarning: 40.0% of the points cannot be placed; you may want to
decrease the size of the markers or use stripplot.
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.12/dist-packages/seaborn/categorical.py:3399:
UserWarning: 25.0% of the points cannot be placed; you may want to
decrease the size of the markers or use stripplot.
warnings.warn(msg, UserWarning)
```



```
sns.stripplot(x=df["Daily_Customer_Count"], y=df["Store_Sales"])
plt.title("Stripplot: Customers vs Sales")
plt.show()
```



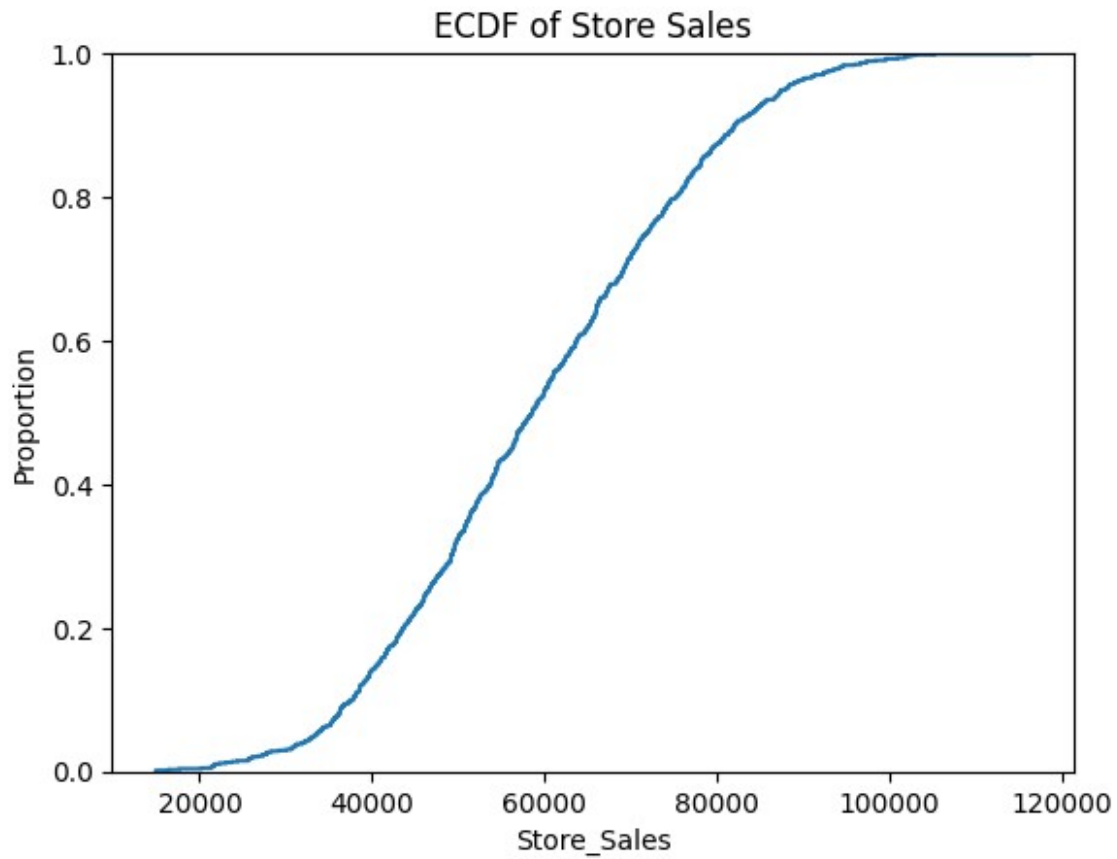
```
df.plot.hexbin(x="Daily_Customer_Count", y="Store_Sales", gridsize=25,  
cmap="Blues")  
plt.title("Hexbin: Customers vs Sales")  
plt.show()
```




```
sns.displot(df["Store_Sales"], kde=True, bins=30)
plt.title("Distribution of Store Sales")
plt.show()
```



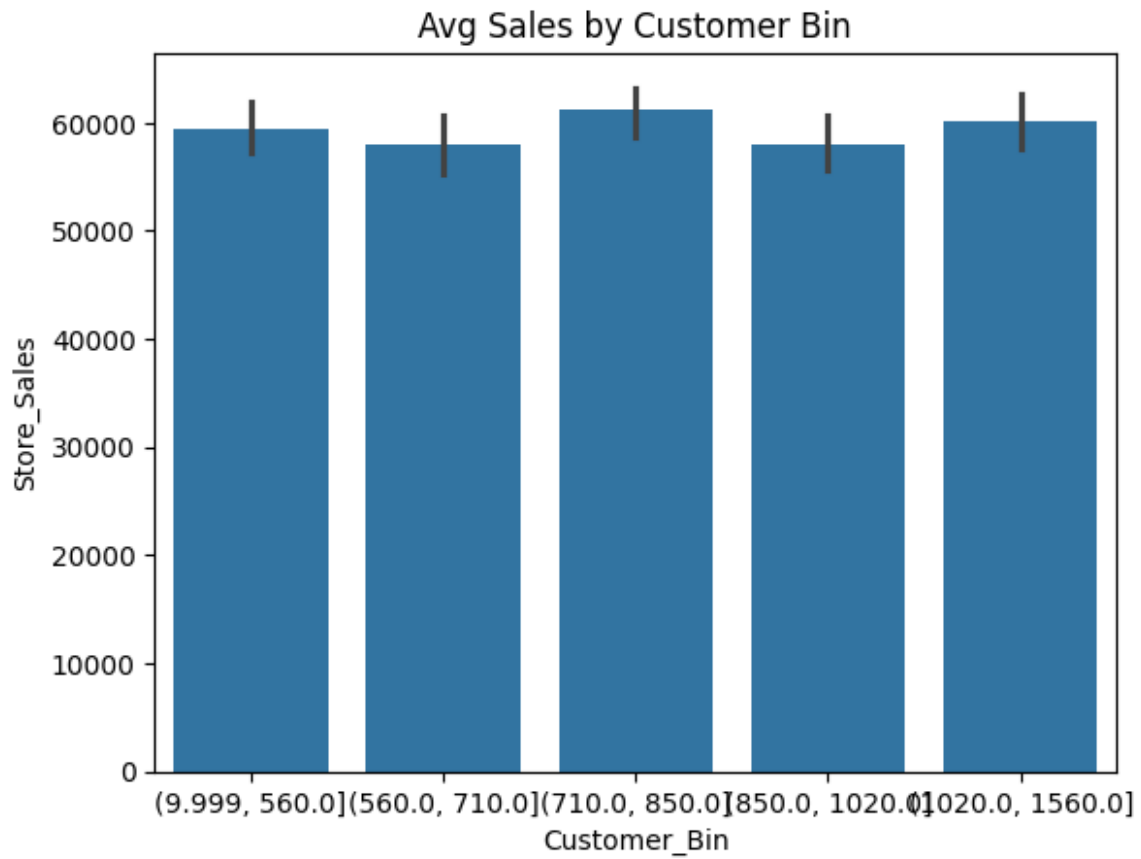
```
sns.ecdfplot(df["Store_Sales"])  
plt.title("ECDF of Store Sales")  
plt.show()
```



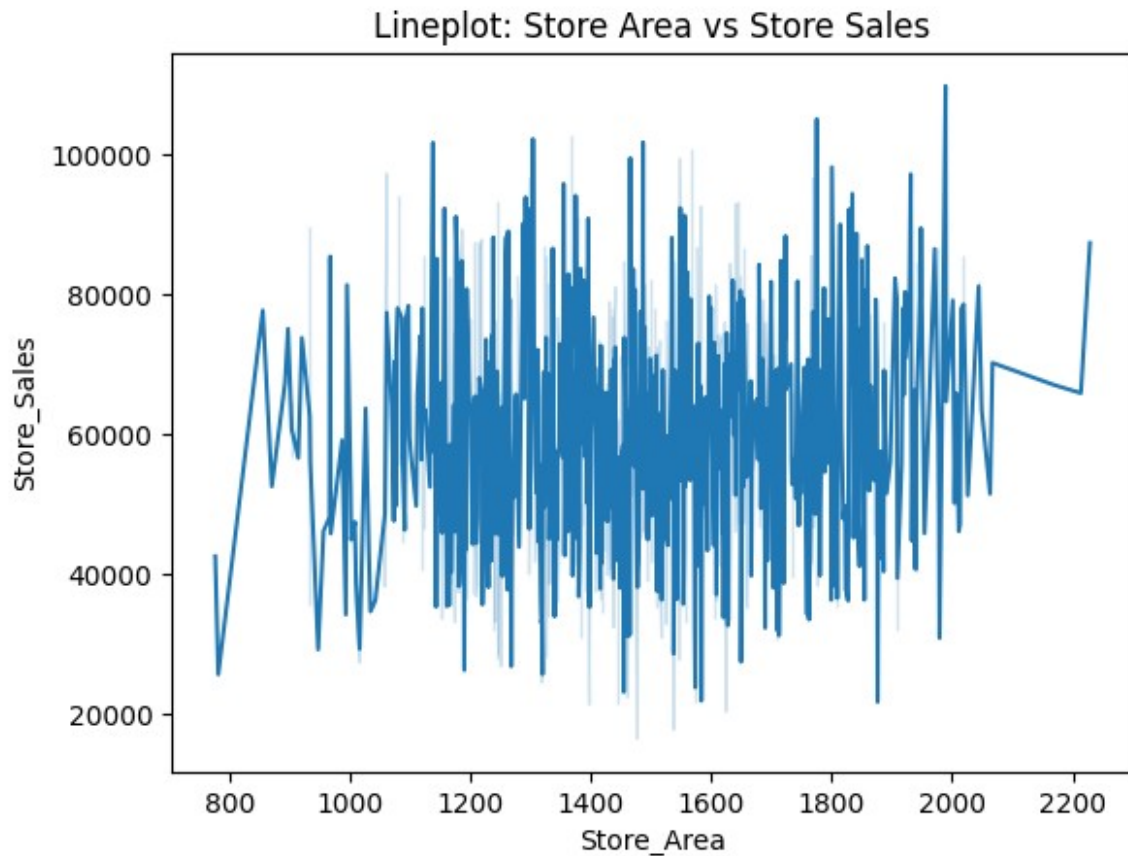
```
df["Customer_Bin"] = pd.qcut(df["Daily_Customer_Count"], q=5)
sns.countplot(x="Customer_Bin", data=df)
plt.title("Countplot of Customer Bins")
plt.show()
```

```
import numpy as np  
  
sns.barplot(x="Customer_Bin", y="Store_Sales", data=df,  
            estimator=np.mean)  
plt.title("Avg Sales by Customer Bin")  
plt.show()
```



```
sns.lineplot(x="Store_Area", y="Store_Sales",  
data=df.sort_values("Store_Area"))  
plt.title("Lineplot: Store Area vs Store Sales")  
plt.show()
```




```
sns.kdeplot(x=df["Daily_Customer_Count"], y=df["Store_Sales"],  
cmap="Reds", shade=True)  
plt.title("Density Plot: Customers vs Sales")  
plt.show()
```

/tmp/ipython-input-1113677589.py:1: FutureWarning:

`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

```
sns.kdeplot(x=df["Daily_Customer_Count"], y=df["Store_Sales"],  
cmap="Reds", shade=True)
```



```
from scipy.stats import f_oneway, chi2_contingency

f_stat, p_val = f_oneway(df['Daily_Customer_Count'],
df['Store_Sales'])
print("ANOVA p-value:", p_val)

ANOVA p-value: 0.0

df['Sales_per_Customer'] = df['Store_Sales'] /
df['Daily_Customer_Count']
df['Items_Density'] = df['Items_Available'] / df['Store_Area']

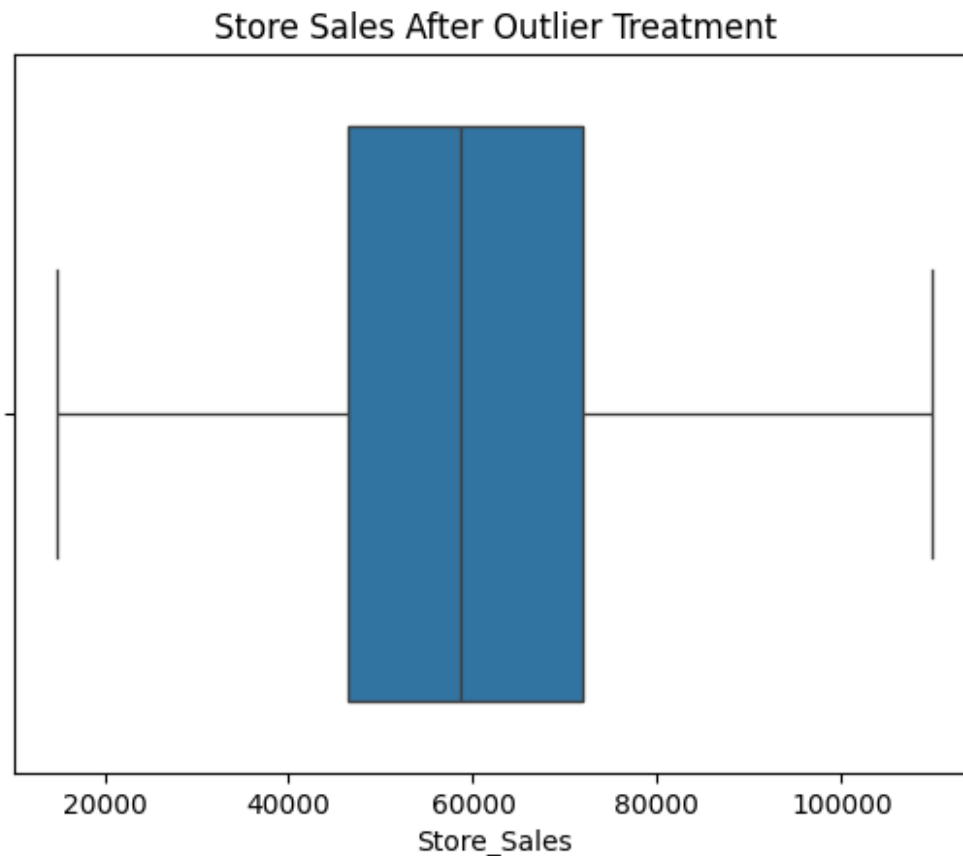
Q1 = df['Store_Sales'].quantile(0.25)
Q3 = df['Store_Sales'].quantile(0.75)
IQR = Q3 - Q1
lower = Q1 - 1.5*IQR
upper = Q3 + 1.5*IQR
df['Store_Sales'] = np.where(df['Store_Sales'] > upper, upper,
df['Store_Sales'])

df['Log_Sales'] = np.log1p(df['Store_Sales'])
```


```
print(df.select_dtypes(include=['int64', 'float64']).corr()  
['Store_Sales'].sort_values(ascending=False))
```

```
Store_Sales      1.000000  
Log_Sales        0.979127  
Sales_per_Customer  0.139473  
Items_Available  0.098097  
Store_Area       0.096759  
Store_ID         0.071254  
Items_Density    0.022556  
Daily_Customer_Count 0.008524  
Name: Store_Sales, dtype: float64
```

```
sns.boxplot(x=df['Store_Sales'])  
plt.title("Store Sales After Outlier Treatment")  
plt.show()
```



```
sns.histplot(df['Log_Sales'], kde=True, bins=30)  
plt.title("Log Transformed Store Sales")  
plt.show()
```



```
print("Final Shape:", df.shape)
print("Final Columns:", df.columns)
```

```
Final Shape: (896, 9)
Final Columns: Index(['Store ID ', 'Store_Area', 'Items_Available',
                      'Daily_Customer_Count',
                      'Store_Sales', 'Customer_Bin', 'Sales_per_Customer',
                      'Items_Density',
                      'Log_Sales'],
                    dtype='object')
```

```
print(df.describe(include="all").T)
```

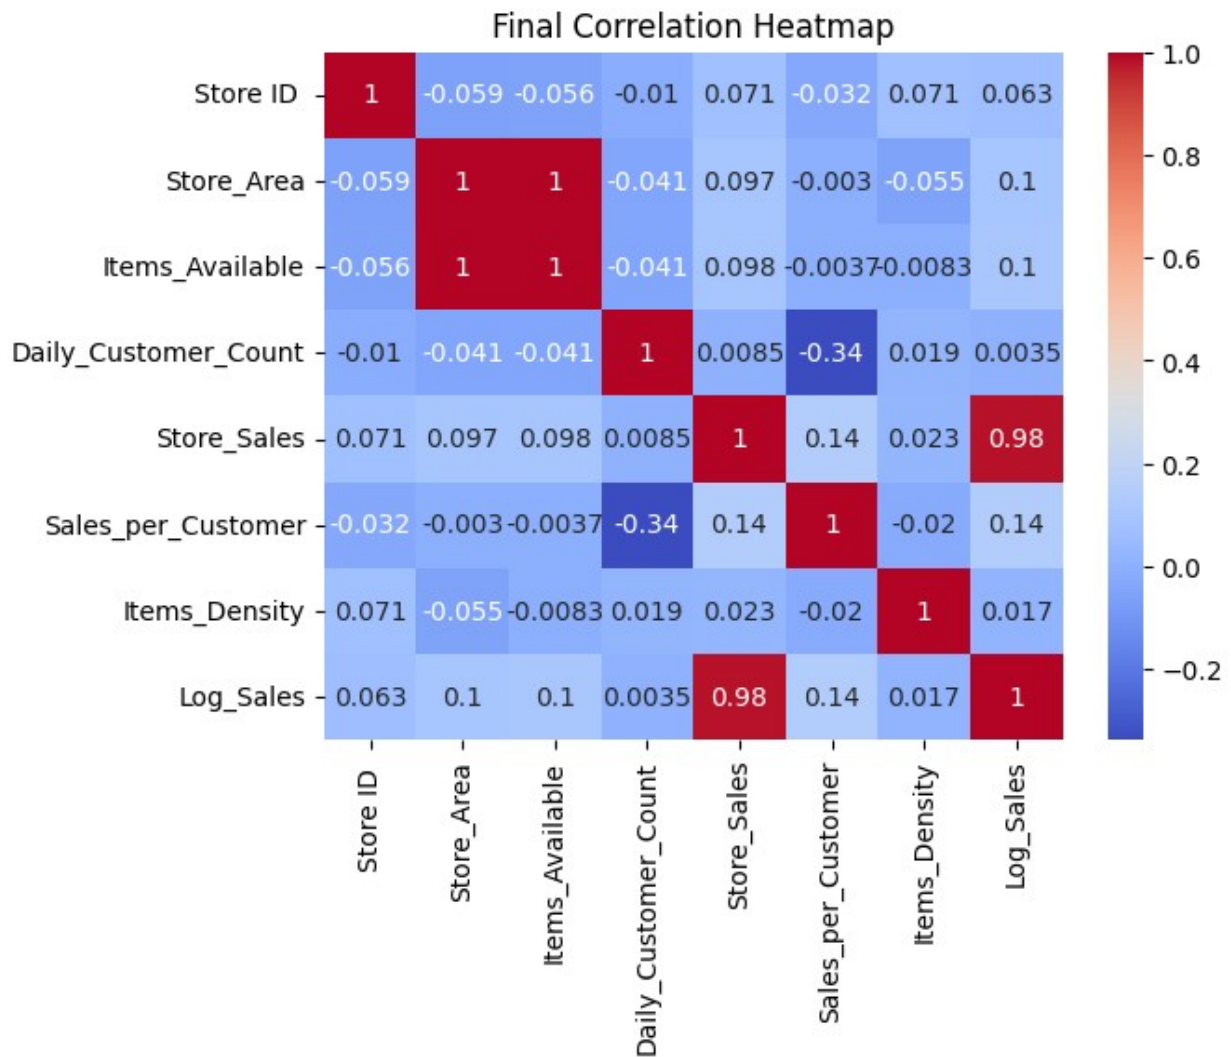
	count	unique	top freq		mean
\					
Store ID	896.0	NaN	NaN	NaN	448.5
Store_Area	896.0	NaN	NaN	NaN	1485.409598
Items_Available	896.0	NaN	NaN	NaN	1782.035714
Daily_Customer_Count	896.0	NaN	NaN	NaN	786.350446
Store_Sales	896.0	NaN	NaN	NaN	59344.125279

Customer_Bin	896	5	(850.0, 1020.0]	184	NaN
Sales_per_Customer	896.0	NaN		NaN	94.034825
Items_Density	896.0	NaN		NaN	1.19978
Log_Sales	896.0	NaN		NaN	10.945311

	std	min	25%	50%	\
Store ID	258.797218	1.0	224.75	448.5	
Store_Area	250.237011	775.0	1316.75	1477.0	
Items_Available	299.872053	932.0	1575.5	1773.5	
Daily_Customer_Count	265.389281	10.0	600.0	780.0	
Store_Sales	17168.248608	14920.0	46530.0	58605.0	
Customer_Bin	NaN	NaN	NaN	NaN	
Sales_per_Customer	162.817175	15.070707	55.875573	75.409239	
Items_Density	0.009409	1.17359	1.19323	1.199695	
Log_Sales	0.313244	9.610525	10.747872	10.978592	

	75%	max
Store ID	672.25	896.0
Store_Area	1653.5	2229.0
Items_Available	1982.75	2667.0
Daily_Customer_Count	970.0	1560.0
Store_Sales	71872.5	109886.25
Customer_Bin	NaN	NaN
Sales_per_Customer	102.853319	4548.0
Items_Density	1.206217	1.22833
Log_Sales	11.182662	11.60721

```
sns.heatmap(df.select_dtypes(include=['int64','float64']).corr(),
annot=True, cmap="coolwarm")
plt.title("Final Correlation Heatmap")
plt.show()
```



```
df.to_csv("Stores_Cleaned.csv", index=False)

from google.colab import files
files.download("Stores_Cleaned.csv")
```